

2-3-4 Trees and B-Trees


Instructors: Subrahmanyam Kalyanasundaram
Karteek Sreenivasaiah

28th September 2020

Multiway search Trees

- ▶ Search trees, but not binary search trees
- ▶ Each node has at least 2 children
- ▶ Each node can store many keys
- ▶ If a node stores d keys, then it has $d + 1$ children
- ▶ All leaf nodes are NIL nodes
- ▶ All leaf nodes are at the same level

Example



All the NIL nodes are not shown above


2-3-4 Trees

- ▶ Multiway search tree where each node has 1, 2 or 3 keys.
- ▶ Consequently, each node has 2, 3 or 4 children
- ▶ What can we say about the height of a 2-3-4 tree?

2-3-4 Trees

- ▶ Multiway search tree where each node has 1, 2 or 3 keys.
- ▶ Consequently, each node has 2, 3 or 4 children
- ▶ What can we say about the height of a 2-3-4 tree?
- ▶ $1/2 \log(n + 1) \leq h \leq \log(n + 1)$

Example



No NIL nodes are shown above

Searching in 2-3-4 tree

- ▶ Similar to BST search
- ▶ Start from the root node
- ▶ Find two keys in the node k_{i-1} and k_i such that the searched value is between these two values
- ▶ Search the subtree between k_{i-1} and k_i
- ▶ Running time?

Searching in 2-3-4 tree

- ▶ Similar to BST search
- ▶ Start from the root node
- ▶ Find two keys in the node k_{i-1} and k_i such that the searched value is between these two values
- ▶ Search the subtree between k_{i-1} and k_i
- ▶ Running time?
- ▶ Takes $O(\log n)$ time

Other query operations

- ▶ How do you find successor/predecessor?

Other query operations

- ▶ How do you find successor/predecessor?
- ▶ How about Max/Min?

Other query operations

- ▶ How do you find successor/predecessor?
- ▶ How about Max/Min?
- ▶ Running time?

Insertion

- ▶ Suppose we want to insert the value x
- ▶ Search for x in the tree
- ▶ If x not found, insert x in the leaf node where it should ideally have been
- ▶ Two cases:

Insertion

- ▶ Suppose we want to insert the value x
- ▶ Search for x in the tree
- ▶ If x not found, insert x in the leaf node where it should ideally have been
- ▶ Two cases:
 - ▶ The node has room for x – it has 1 or 2 values only
 - ▶ The node is full – it has already 3 values

INSERT(x)

Case 1

The node has room for x

INSERT(x)

Case 1

The node has room for x

Resolution:

- ▶ We simply add x to the leaf node where it should have been
- ▶ Maintain the keys in sorted order

Case 1

17

15		
----	--	--

Case 1

15	17	
----	----	--

Case 1

16

15	17	
----	----	--

Case 1

15	16	17
----	----	----

INSERT(x)

Case 2

The node has no room for x

INSERT(x)

Case 2

The node has no room for x

Resolution:

- ▶ Adding x to the node results in 4 keys
- ▶ We cannot have 4 keys

INSERT(x)


Case 2

The node has no room for x


Resolution:

- ▶ Adding x to the node results in 4 keys
- ▶ We cannot have 4 keys
- ▶ We split the node and promote the median


Case 2



Case 2




Case 2




Split and promote!

Case 2




Case 2



- ▶ Can we promote any other key?
- ▶ What if the parent node doesn't have room?

Case 2



- ▶ Can we promote any other key? The other median.
- ▶ What if the parent node doesn't have room? Recurse up!

INSERT Example

On the board

Deletion

- ▶ We want to insert the value x
- ▶ If x is in the leaf, we delete x from the leaf
- ▶ Else, we swap x with its successor/predecessor and delete the succ/pred

Deletion

- ▶ We want to insert the value x
- ▶ If x is in the leaf, we delete x from the leaf
- ▶ Else, we swap x with its successor/predecessor and delete the succ/pred
- ▶ **Note:** The succ/pred will always be in a leaf node if x is not in a leaf.

Deletion

- ▶ We want to insert the value x
 - ▶ If x is in the leaf, we delete x from the leaf
 - ▶ Else, we swap x with its successor/predecessor and delete the succ/pred
 - ▶ **Note:** The succ/pred will always be in a leaf node if x is not in a leaf.
-
- ▶ From now on, we discuss deletion from leaf node

Deletion

Cases:

Deletion

Cases:

- ▶ The node has another key apart from x
- ▶ x is the only value in the node, but can “borrow” from sibling
- ▶ x is the only value in the node and cannot “borrow” from sibling

DELETE(x)

Case 1

The node has another key

DELETE(x)

Case 1

The node has another key

Resolution:

- We simply remove the key x

Case 1

15	16	17
----	----	----

- Delete 17

Case 1

15	16	
----	----	--

► Delete 17 Done!

Case 1

15	16	
----	----	--

- ▶ Delete 17 Done!
- ▶ Delete 16

Case 1

15		
----	--	--

- ▶ Delete 17 Done!
- ▶ Delete 16 Done!

Case 1

15		
----	--	--

- ▶ Delete 17 Done!
- ▶ Delete 16 Done!
- ▶ Delete 15? Next Cases!

DELETE(x)

Case 2

The node only one key, x

Can “borrow” from sibling node

DELETE(x)


Case 2

The node only one key, x
Can “borrow” from sibling node

Resolution:


- ▶ Adjacent sibling must have ≥ 2 keys
- ▶ Can borrow from the adjacent sibling, through the parent

Case 2




► Delete 15

Case 2




- Delete 15

Case 2




- ▶ Delete 15
- ▶ 13 is transferred to the parent node, and 14 is brought down
- ▶ Similar to

Case 2



- ▶ Delete 15
- ▶ 13 is transferred to the parent node, and 14 is brought down
- ▶ Similar to **Rotation!**
- ▶ Like in rotation, we transfer one child of the sibling node

Case 2



- ▶ Delete 15
- ▶ 13 is transferred to the parent node, and 14 is brought down
- ▶ Similar to **Rotation!**
- ▶ Like in rotation, we transfer one child of the sibling node
- ▶ What if we cannot borrow from sibling?

DELETE(x)

Case 3

The node only one key, x
Cannot borrow from sibling

DELETE(x)

Case 3

The node only one key, x
Cannot borrow from sibling

Resolution:

- Merge with a sibling

DELETE(x)


Case 3

The node only one key, x
Cannot borrow from sibling

Resolution:


- ▶ Merge with a sibling
- ▶ Need to bring a key down from parent node

Case 3




- ▶ Delete 14
- ▶ Cannot borrow from either sibling

Case 3




- ▶ Delete 14
- ▶ Cannot borrow from either sibling
- ▶ Once we remove the node, we have an issue

Case 3




- ▶ Delete 14
- ▶ Cannot borrow from either sibling
- ▶ Once we remove the node, we have an issue
- ▶ Bring down a key from parent

Case 3



- ▶ Delete 14
- ▶ Cannot borrow from either sibling
- ▶ Once we remove the node, we have an issue
- ▶ Bring down a key from parent
- ▶ What if parent has only one key?

Case 3



- ▶ Delete 14
- ▶ Cannot borrow from either sibling
- ▶ Once we remove the node, we have an issue
- ▶ Bring down a key from parent
- ▶ What if parent has only one key? **Recurse!**

DELETE Example

On the board

Summary of INSERT and DELETE

- ▶ At each node, we do an $O(1)$ time operation
 - ▶ Add/remove key
 - ▶ Split/Merge
 - ▶ Borrow from sibling
 - ▶ Promote to/bring down from parent
- ▶ We may go up the tree as well, upto height h
- ▶ Running time is $O(h) = O(\log n)$

Questions

- ▶ Think about how the insert/delete operations compare with the operations in Red-Black Trees.
- ▶ Could we extend this notion to an (a, b) -tree? What conditions should be satisfied by a and b ?

2-3-4 Trees

- ▶ What can we say about the height of a 2-3-4 tree?
- ▶ $1/2 \log(n+1) \leq h \leq \log(n+1)$
- ▶ All operations, query and modify, are $O(\log n)$

2-3-4 Trees: Implementation

Each node contains:

- ▶ d , the number of keys in the node
- ▶ x_1, x_2, \dots, x_d , the keys in increasing order
- ▶ The pointers to the $d + 1$ children
- ▶ A bit that indicates whether the node is an external node

(a, b) -tree

- ▶ 2-3-4 Trees are $(2, 4)$ -trees
- ▶ In (a, b) -trees, each node has at least $a - 1$ and at most $b - 1$ keys
- ▶ So each node has at least a and at most b children
- ▶ The lower bound of $a - 1$ is **not** applicable to the root

(a, b) -tree

- ▶ 2-3-4 Trees are $(2, 4)$ -trees
- ▶ In (a, b) -trees, each node has at least $a - 1$ and at most $b - 1$ keys
- ▶ So each node has at least a and at most b children
- ▶ The lower bound of $a - 1$ is **not** applicable to the root
- ▶ Exercise: Show that $2(a - 1) \leq b - 1$ has to be satisfied.

B-Tree

- ▶ B-Trees are (a, b) -trees with large values of a and b
- ▶ In a large database, the tree may be stored in the secondary memory
- ▶ Accessing a “page” takes time
- ▶ It helps if the entire page is a node

INSERT

- ▶ Search for the key, insert at leaf, may need to recurse up
- ▶ The procedure for INSERT in CLRS avoids recursing up
- ▶ This procedure splits every full node pre-emptively while searching
- ▶ If the leaf node needs to be split, then the parent is sure to have room to accommodate the median

INSERT

- ▶ Search for the key, insert at leaf, may need to recurse up
- ▶ The procedure for INSERT in CLRS avoids recursing up
- ▶ This procedure splits every full node pre-emptively while searching
- ▶ If the leaf node needs to be split, then the parent is sure to have room to accommodate the median
- ▶ Avoids the upwards recursion!

DELETE

- ▶ This also can be executed in one-pass
- ▶ During the search for the node, pre-emptively merge the nodes that have min no. of elements, and can't borrow

DELETE

- ▶ This also can be executed in one-pass
- ▶ During the search for the node, pre-emptively merge the nodes that have min no. of elements, and can't borrow
- ▶ Exercise: Read the procedure in CLRS